



AKASH-M-hub / EXNO-4-DS



&lt;&gt; Code



Pull requests



Actions



Projects



Wiki



Security



Insights



Settings



EXNO-4-DS /

README.md

in main

Cancel changes

Commit changes...

Edit

Preview

☐ Show Diff

# EXNO:4-DS

## AIM:

To read the given data and perform Feature Scaling and Feature Selection process and save the data to a file.

## ALGORITHM:

STEP 1:Read the given Data. STEP 2:Clean the Data Set using Data Cleaning Process. STEP 3:Apply Feature Scaling for the feature in the data set. STEP 4:Apply Feature Selection for the feature in the data set. STEP 5:Save the data to the file.

## FEATURE SCALING:

1. Standard Scaler: It is also called Z-score normalization. It calculates the z-score of each value and replaces the value with the calculated Z-score. The features are then rescaled with  $\bar{x} = 0$  and  $\sigma = 1$
2. MinMaxScaler: It is also referred to as Normalization. The features are scaled between 0 and 1. Here, the mean value remains same as in Standardization, that is, 0.
3. Maximum absolute scaling: Maximum absolute scaling scales the data to its maximum value; that is, it divides every observation by the maximum value of the variable. The result of the preceding transformation is a distribution in which the values vary approximately within the range of -1 to 1.
4. RobustScaler: RobustScaler transforms the feature vector by subtracting the median and then dividing by the interquartile range (75% value — 25% value).

## FEATURE SELECTION:

---

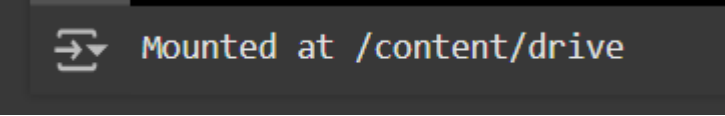
Feature selection is to find the best set of features that allows one to build useful models. Selecting the best features helps the model to perform well. The feature selection techniques used are: 1.Filter Method 2.Wrapper Method 3.Embedded Method

## CODING AND OUTPUT:

---

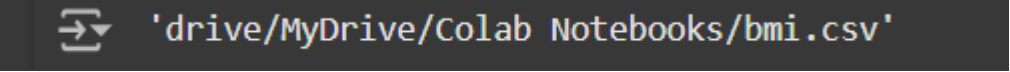
```
from google.colab import drive
```

```
drive.mount('/content/drive')
```

A dark-themed terminal window showing the output of the drive.mount() command. It displays a double arrow icon followed by the text "Mounted at /content/drive".

```
Mounted at /content/drive
```

```
ls drive/MyDrive/'Colab Notebooks'/bmi.csv
```

A dark-themed terminal window showing the output of the ls command. It displays a double arrow icon followed by the file path "drive/MyDrive/Colab Notebooks/bmi.csv".

```
'drive/MyDrive/Colab Notebooks/bmi.csv'
```

```
ls drive/MyDrive/'Colab Notebooks'/'income(1) (1).csv'
```

```
⇌ 'drive/MyDrive/Colab Notebooks/income(1) (1).csv'
```


```
import pandas as pd
```

```
import numpy as np
```

```
from scipy import stats
```

```
df=pd.read_csv("drive/MyDrive/Colab Notebooks/bmi.csv")
```


```
df
```



	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...	...	...	...	...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5


500 rows × 4 columns

```
df.head()
```



	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3

```
df.dropna()
```



	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...	...	...	...	...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

500 rows × 4 columns


```
max_vals = np.max(np.abs(df[['Height','Weight']]))
```

```
max_vals
```

```
max_vals
```

```
df1=pd.read_csv("drive/MyDrive/Colab Notebooks/bmi.csv")
```

df1



	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...	...	...	...	...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

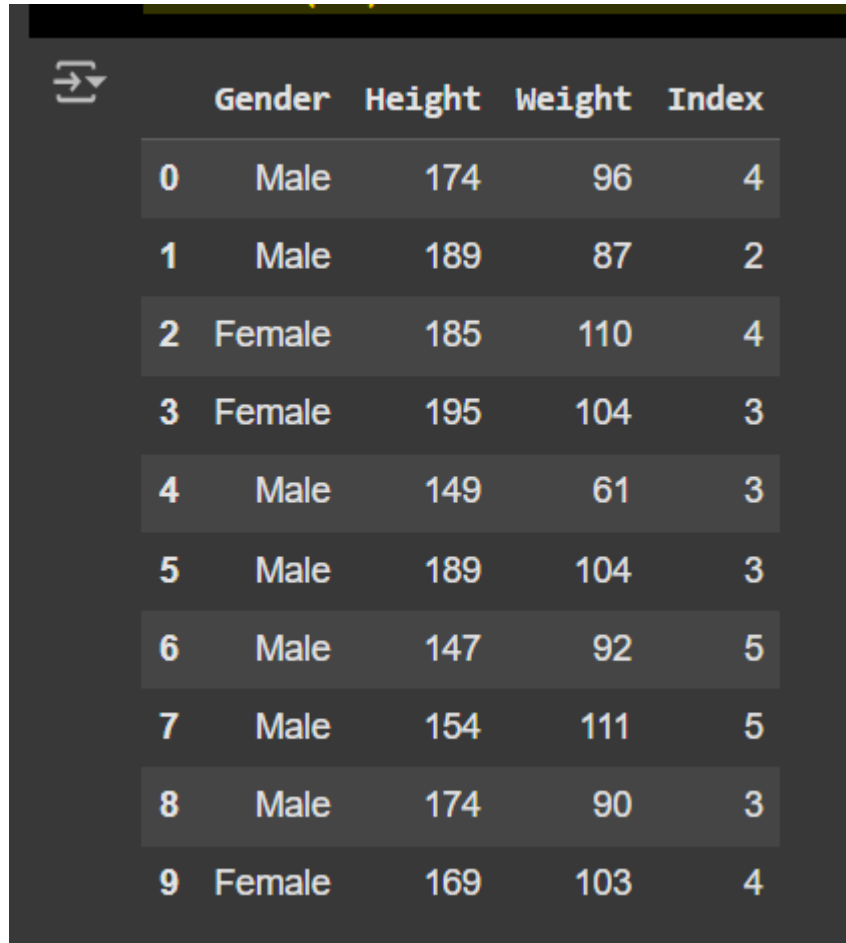
500 rows × 4 columns

```
from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
```

```
df1[['Height','Weight']] = sc.fit_transform(df1[['Height','Weight']])
```

```
df.head(10)
```



	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
5	Male	189	104	3
6	Male	147	92	5
7	Male	154	111	5
8	Male	174	90	3
9	Female	169	103	4

```
from sklearn.preprocessing import MinMaxScaler
```

```
Scaler=MinMaxScaler()
```

```
df[['Height','Weight']]=Scaler.fit_transform(df[['Height','Weight']])
```

```
df.head(0)
```





**Gender Height Weight Index**

```
df2=pd.read_csv("drive/MyDrive/Colab Notebooks/bmi.csv")
```

df2



**Gender Height Weight Index**

**0 Male 174 96 4**

**1 Male 189 87 2**

**2 Female 185 110 4**

**3 Female 195 104 3**

**4 Male 149 61 3**

**... ... ... ...**

**495 Female 150 153 5**

**496 Female 184 121 4**

**497 Female 141 136 5**

**498 Male 150 95 5**

**499 Male 173 131 5**


**500 rows × 4 columns**

```
from sklearn.preprocessing import Normalizer
```

```
Scaler=Normalizer()
```

```
df2[['Height','Weight']]=Scaler.fit_transform(df2[['Height','Weight']])
```

```
df2
```




	Gender	Height	Weight	Index
0	Male	0.875578	0.483077	4
1	Male	0.908381	0.418144	2
2	Female	0.859536	0.511075	4
3	Female	0.882353	0.470588	3
4	Male	0.925448	0.378875	3
...	...	...	...	...
495	Female	0.700071	0.714073	5
496	Female	0.835527	0.549450	4
497	Female	0.719753	0.694230	5
498	Male	0.844819	0.535052	5
499	Male	0.797227	0.603680	5

500 rows × 4 columns

```
df3=pd.read_csv("drive/MyDrive/Colab Notebooks/bmi.csv")
```

```
df3
```



	Gender	Height	Weight	Index
0	Male	174	96	4
1	Male	189	87	2
2	Female	185	110	4
3	Female	195	104	3
4	Male	149	61	3
...	...	...	...	...
495	Female	150	153	5
496	Female	184	121	4
497	Female	141	136	5
498	Male	150	95	5
499	Male	173	131	5

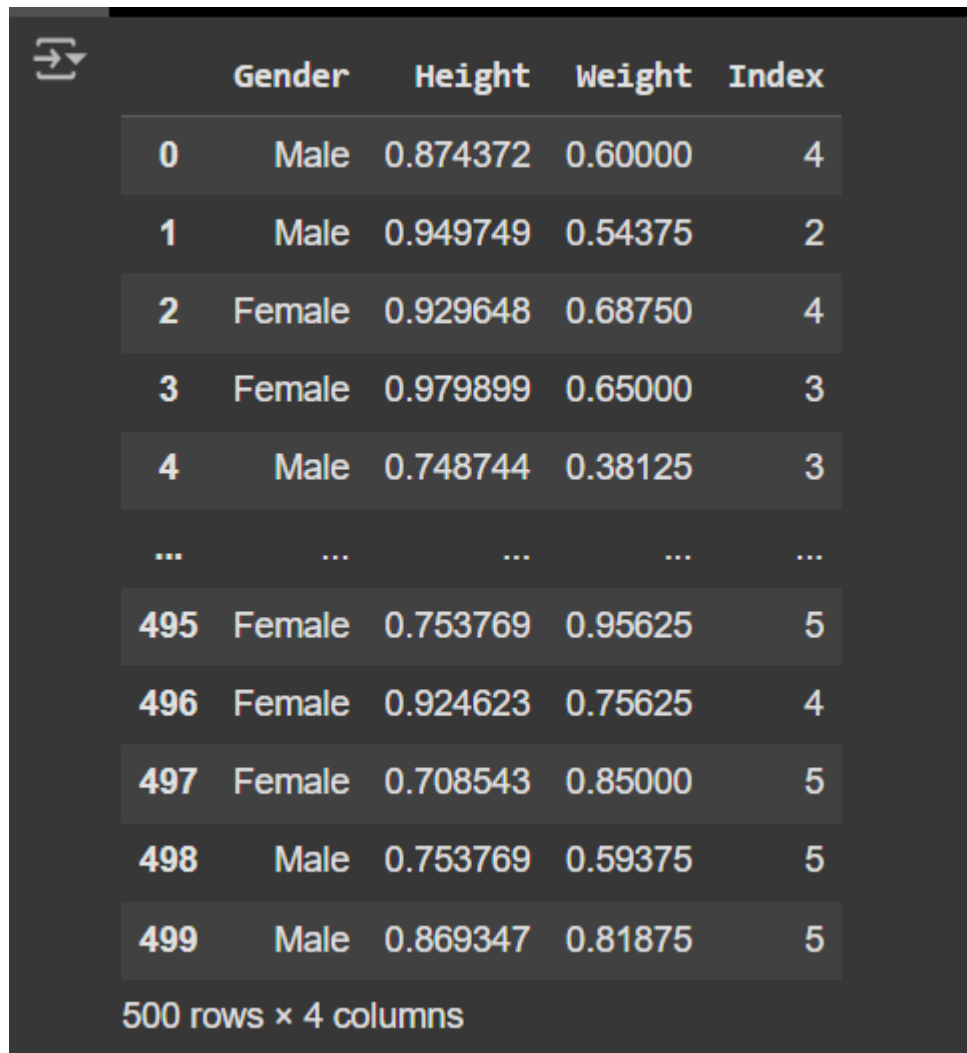
500 rows × 4 columns

```
from sklearn.preprocessing import MaxAbsScaler
```

```
Scaler=MaxAbsScaler()
```

```
df3[['Height','Weight']]=Scaler.fit_transform(df3[['Height','Weight']])
```

df3



	Gender	Height	Weight	Index
0	Male	0.874372	0.60000	4
1	Male	0.949749	0.54375	2
2	Female	0.929648	0.68750	4
3	Female	0.979899	0.65000	3
4	Male	0.748744	0.38125	3
...	...	...	...	...
495	Female	0.753769	0.95625	5
496	Female	0.924623	0.75625	4
497	Female	0.708543	0.85000	5
498	Male	0.753769	0.59375	5
499	Male	0.869347	0.81875	5

500 rows x 4 columns

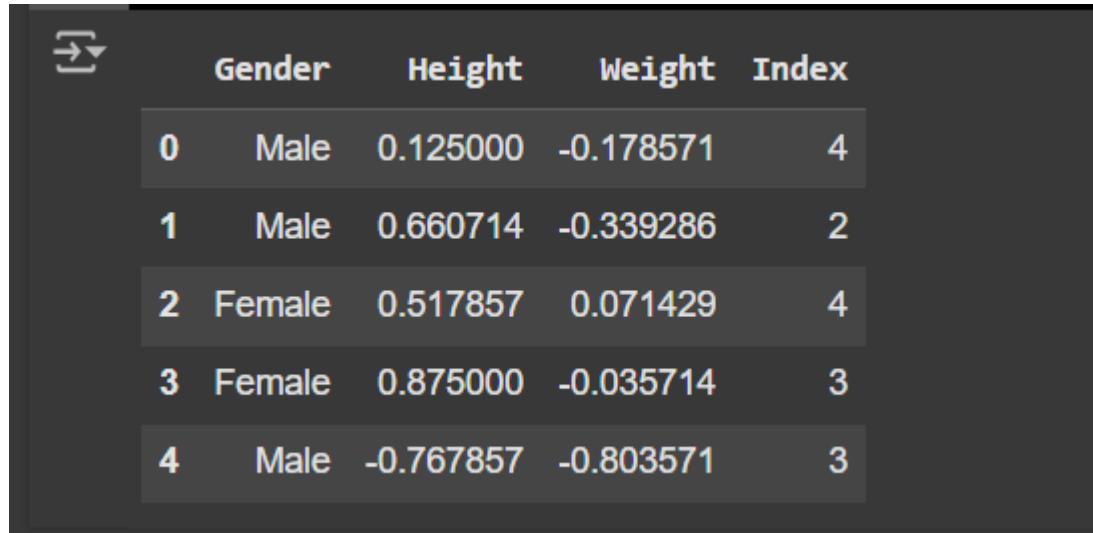
```
df4=pd.read_csv("drive/MyDrive/Colab Notebooks/bmi.csv")
```

```
from sklearn.preprocessing import RobustScaler
```

```
Scaler=RobustScaler()
```

```
df4[['Height','Weight']]=Scaler.fit_transform(df4[['Height','Weight']])
```

```
df4.head()
```

A screenshot of a Jupyter Notebook interface. On the left, there is a sidebar with a file explorer icon. The main area displays the output of a pandas DataFrame head() function. The DataFrame has five columns: 'Index', 'Gender', 'Height', 'Weight', and 'Index'. The data is as follows:

	Gender	Height	Weight	Index
0	Male	0.125000	-0.178571	4
1	Male	0.660714	-0.339286	2
2	Female	0.517857	0.071429	4
3	Female	0.875000	-0.035714	3
4	Male	-0.767857	-0.803571	3

```
import matplotlib
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
import statsmodels.api as sm
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.feature_selection import RFE
```

```
from sklearn.linear_model import RidgeCV,LassoCV,Ridge,Lasso
```

```
from sklearn.feature_selection import SelectKBest

from sklearn.feature_selection import mutual_info_classif

from sklearn.feature_selection import mutual_info_regression

from sklearn.feature_selection import chi2

df=pd.read_csv('drive/MyDrive/Colab Notebooks/income(1) (1).csv')

df.columns
```

```
⇒ Index(['age', 'JobType', 'EdType', 'maritalstatus', 'occupation',
        'relationship', 'race', 'gender', 'capitalgain', 'capitalloss',
        'hoursperweek', 'nativecountry', 'SalStat'],
        dtype='object')
```

```
df1.columns
```

```
⇒ Index(['Gender', 'Height', 'Weight', 'Index'], dtype='object')
```

```
import pandas as pd

from sklearn.feature_selection import SelectKBest


from sklearn.feature_selection import chi2

data=pd.read_csv('drive/MyDrive/Colab Notebooks/bmi.csv')

data=data.dropna()

df.columns
```

df



	age	JobType	EdType	maritalstatus	occupation	relationship	race	gender	capitalgain	capitalloss	hoursperweek	nativecountry	SalStat
0	45	Private	HS-grad	Divorced	Adm-clerical	Not-in-family	White	Female	0	0	28	United-States	less than or equal to 50,000
1	24	Federal-gov	HS-grad	Never-married	Armed-Forces	Own-child	White	Male	0	0	40	United-States	less than or equal to 50,000
2	44	Private	Some-college	Married-civ-spouse	Prof-specialty	Husband	White	Male	0	0	40	United-States	greater than 50,000
3	27	Private	9th	Never-married	Craft-repair	Other-relative	White	Male	0	0	40	Mexico	less than or equal to 50,000
4	20	Private	Some-college	Never-married	Sales	Not-in-family	White	Male	0	0	35	United-States	less than or equal to 50,000
...	...	...	...	...	...	...	...	...	...	...	...	...	...
31973	34	Local-gov	HS-grad	Never-married	Farming-fishing	Not-in-family	Black	Male	594	0	60	United-States	less than or equal to 50,000
31974	34	Local-gov	Some-college	Never-married	Protective-serv	Not-in-family	White	Female	0	0	40	United-States	less than or equal to 50,000
31975	23	Private	Some-college	Married-civ-spouse	Adm-clerical	Husband	White	Male	0	0	40	United-States	less than or equal to 50,000
31976	42	Local-gov	Some-college	Married-civ-spouse	Adm-clerical	Wife	White	Female	0	0	40	United-States	less than or equal to 50,000
31977	29	Private	Bachelors	Never-married	Prof-specialty	Not-in-family	White	Male	0	0	40	United-States	less than or equal to 50,000

31978 rows × 13 columns

```
import pandas as pd
```


```
import numpy as np
```

```
from scipy.stats import chi2_contingency
```

```
import seaborn as sns
```

```
tips=sns.load_dataset('tips')
```

```
tips.head()
```



	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	Female	No	Sun	Dinner	2
1	10.34	1.66	Male	No	Sun	Dinner	3
2	21.01	3.50	Male	No	Sun	Dinner	3
3	23.68	3.31	Male	No	Sun	Dinner	2
4	24.59	3.61	Female	No	Sun	Dinner	4

## RESULT:

---

THUS THE ABOVE CODE IS EXECUTED SUCCESSFULLY

